

**Full Stack AI Software Development**

**Advanced Database Design  
and Development**

**Job Connector Program**

# Outline

## **Database Design and Normalization**

Learn how to structure databases efficiently using normalization to reduce redundancy and improve consistency.

## **Database Key, JOIN and Relation**

Understand primary and foreign keys, table relationships, and how to combine data using SQL JOINS.

## **Database Aggregate Function**

Explore SQL aggregate functions like COUNT, SUM, AVG, MIN, and MAX to analyze and summarize data.

## **SubQuery, Transaction and Indexing**

Learn how to use subqueries, manage transactions, and optimize performance with indexing techniques.

# Database Design

Database design can be generally defined as a collection of tasks or processes that enhance the designing, development, implementation, and maintenance of enterprise data management system. Designing a proper database reduces the maintenance cost thereby improving data consistency and the cost-effective measures are greatly influenced in terms of disk storage space. Therefore, there has to be a brilliant concept of designing a database. The designer should follow the constraints and decide how the elements correlate and what kind of data must be stored.

# What is good database design?

A good database design is, therefore, one that:

- Divides your information into subject-based tables to reduce redundant data.
- Provides Access with the information it requires to join the information in the tables together as needed.
- Helps support and ensure the accuracy and integrity of your information.
- Accommodates your data processing and reporting needs.

# The Design Process

The design process consists of the following steps:

- **Determine the purpose of your database** - This helps prepare you for the remaining steps.
- **Find and organize the information required** - Gather all of the types of information you might want to record in the database, such as product name and order number.
- **Divide the information into tables** - Divide your information items into major entities or subjects, such as Products or Orders. Each subject then becomes a table.
- **Turn information items into columns** - Decide what information you want to store in each table. Each item becomes a field, and is displayed as a column in the table. For example, an Employees table might include fields such as Last Name and Hire Date.

# The Design Process Continue

- **Specify primary keys** - Choose each table's primary key. The primary key is a column that is used to uniquely identify each row. An example might be Product ID or Order ID.
- **Set up the table relationships** - Look at each table and decide how the data in one table is related to the data in other tables. Add fields to tables or create new tables to clarify the relationships, as necessary.
- **Refine your design** - Analyze your design for errors. Create the tables and add a few records of sample data. See if you can get the results you want from your tables. Make adjustments to the design, as needed.
- **Apply the normalization rules** - Apply the data normalization rules to see if your tables are structured correctly. Make adjustments to the tables, as needed.

# Database Normalization


The elementary concepts used in database normalization are:

- **Keys.** Column attributes that identify a database record uniquely.
- **Functional Dependencies.** Constraints between two attributes in a relation.
- **Normal Forms.** Steps to accomplish a certain quality of a database.

# Database Normal Form

Normalizing a database is achieved through a set of rules known as normal forms. The central concept is to help a database designer achieve the desired quality of a relational database.

A database is normalized when it fulfills the third normal form. Further steps in normalization make the database design complicated and could compromise the functionality of the system.



# Stage of Normalization

Stage	Redundancy Anomalies Addressed
Unnormalized Form (UNF)	The state before any normalization. Redundant and complex values are present.
First Normal Form (1NF)	Repeating and complex values split up, making all instances atomic.
Second Normal Form (2NF)	Partial dependencies decompose to new tables. All rows functionally depend on the primary key.
Third Normal Form (3NF)	Transitive dependencies decompose to new tables. Non-key attributes depend on the primary key.
Boyce-Codd Normal Form (BCNF)	Transitive and partial functional dependencies for all candidate keys decompose to new tables.
Fourth Normal Form (4NF)	Removal of multivalued dependencies.
Fifth Normal Form (5NF)	Removal of <b>JOIN</b> dependencies.

# Database Key

What is a KEY?

A database key is an attribute or a group of features that uniquely describes an entity in a table. The types of keys used in normalization are:

- **Super Key.** A set of features that uniquely define each record in a table.
- **Candidate Key.** Keys selected from the set of super keys where the number of fields is minimal.
- **Primary Key.** The most appropriate choice from the set of candidate keys serves as the table's primary key.
- **Foreign Key.** The primary key of another table.
- **Composite Key.** Two or more attributes together form a unique key but are not keys individually.

# Database Key - Super Key

Some examples of super keys in the table are:

- employeeID
- (employeeID, name)
- email

All super keys can serve as a unique identifier for each row. On the other hand, the employee's name or age are not unique identifiers because two people could have the same name or age.

employeeID	name	age	email
1	Adam A.	30	adam.a@email.com
2	Jacob J.	27	jacob.j@email.com
3	David D.	35	david.d@email.com

# Database Key - Candidate Key

The candidate keys come from the set of super keys where the number of fields is minimal. The choice comes down to two options:

- employeeID
- email

Both options contain a minimal number of fields, making them optimal candidate keys.

employeeID	name	age	email
1	Adam A.	30	adam.a@email.com
2	Jacob J.	27	jacob.j@email.com
3	David D.	35	david.d@email.com

# Database Key - Primary Key

- The Primary keys constraint uniquely identifies each record in a database table.
- Primary keys must contain UNIQUE values, and cannot contain NULL values.
- A table can have only one primary key, which may consist of single or multiple fields. Example: No\_KTP, product\_key, ID\_sidikjari

# Database Key - Foreign Key

A foreign key is a field or a column that is used to establish a link between two tables. In simple words you can say that, a foreign key in one table used to point primary key in another table.

# Foreign Key

The "S\_Id" column in the 1st table is the PRIMARY KEY in the 1st table.

The "S\_Id" column in the 2nd table is a FOREIGN KEY in the 2nd table.

S_Id	LastName	FirstName	CITY
1	MAURYA	AJEET	ALLAHABAD
2	JAISWAL	RATAN	GHAZIABAD
3	ARORA	SAUMYA	MODINAGAR

O_Id	OrderNo	S_Id
1	99586465	2
2	78466588	2
3	22354846	3

# Join Table

A **JOIN** clause is used to combine rows from two or more tables, based on a related column between them.

Notice that the "CustomerID" column in the "Orders" table refers to the "CustomerID" in the "Customers" table. The relationship between the two tables above is the "CustomerID" column.

OrderID	CustomerID	OrderDate
10308	2	1996-09-18
10309	37	1996-09-19
10310	77	1996-09-20

CustomerID	CustomerName	ContactName	Country
1	Alfreds Futterkiste	Maria Anders	Germany
2	Ana Trujillo Emparedados y helados	Ana Trujillo	Mexico
3	Antonio Moreno Taquería	Antonio Moreno	Mexico

# Join Table Syntax



```
SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate  
FROM Orders  
INNER JOIN Customers ON Orders.CustomerID = Customer.CustomerID;
```

In this part we would like to join the tables between Orders with Customers table. The Orders refer to a table name, while the OrderID refers to a column name inside the Orders table. The Customers refer to a table name, while the CustomerName refers to a column name inside the Customers table.

# Join Table Syntax



```
SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate  
FROM Orders  
INNER JOIN Customers ON Orders.CustomerID = Customer.CustomerID;
```

“FROM Orders” this mean source of the table, while “INNER JOIN” represent to type of join syntax and following by table name “INNER JOIN Customers” this mean that Customers table would be join with INNER JOIN type to table Orders.


# Join Table Syntax



```
SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate  
FROM Orders  
INNER JOIN Customers ON Orders.CustomerID = Customer.CustomerID;
```

“ON” represent to highlight which column that could be relate one table to another using the **foreign key**. In this case CustomerID column on table Orders is the foreign key.

# Type of SQL Join




Go to <https://dev.mysql.com/doc/index-other.html> and download sakila db. Extract and import sakila data into your MySQL.

We would use this sample data to learn about join

# Inner Join

The INNER JOIN keyword selects records that have matching values in both tables.


```
SELECT column_name  
FROM table1  
INNER JOIN table2  
ON table1.column_name = table2.column_name;
```



# Inner Join Example

The INNER JOIN keyword selects records that have matching values in both tables.

```
SELECT *
FROM country
INNER JOIN city
ON country.country_id = city.country_id;
```




# Left Join

The LEFT JOIN keyword returns all records from the left table (table1), and the matching records from the right table (table2). The result is 0 records from the right side, if there is no match.

```
SELECT column_name  
FROM table1  
LEFT JOIN table2  
ON table1.column_name = table2.column_name;
```

LEFT JOIN




# Left Join Example

The LEFT JOIN keyword returns all records from the left table (table1), and the matching records from the right table (table2). The result is 0 records from the right side, if there is no match.

```
SELECT *  
FROM customer  
LEFT JOIN actor  
ON customer.last_name = actor.last_name;
```


LEFT JOIN



# Right Join

The RIGHT JOIN keyword returns all records from the right table (table2), and the matching records from the left table (table1). The result is 0 records from the left side, if there is no match.


```
SELECT column_name  
FROM table1  
RIGHT JOIN table2  
ON table1.column_name = table2.column_name;
```



# Right Join Example

The RIGHT JOIN keyword returns all records from the right table (table2), and the matching records from the left table (table1). The result is 0 records from the left side, if there is no match.

```
SELECT *  
FROM customer  
RIGHT JOIN actor  
ON customer.last_name = actor.last_name;
```



# Cross Join

The CROSS JOIN keyword returns all records from both tables (table1 and table2).

**CROSS JOIN can potentially return very large result-sets!**


The CROSS JOIN keyword returns all matching records from both tables whether the other table matches or not, those rows will be listed as well.

If you add a WHERE clause (if table1 and table2 has a relationship), the CROSS JOIN will produce the same result as the INNER JOIN clause:



```
SELECT column_name  
FROM table1  
CROSS JOIN table2;
```

**CROSSJOIN**



# Cross Join


The CROSS JOIN keyword returns all records from both tables (table1 and table2).

**CROSS JOIN can potentially return very large result-sets!**

The CROSS JOIN keyword returns all matching records from both tables whether the other table matches or not, those rows will be listed as well.


If you add a WHERE clause (if table1 and table2 has a relationship), the CROSS JOIN will produce the same result as the INNER JOIN clause:

```
SELECT * FROM customer  
CROSS JOIN actor;
```



# Database Relationships in MySQL

Database relationship means how the data in one table is related to the data in another table. In RDBMS (Relational Database Management System). The term “Relational” refers to the tables with Relations. Relationships between two tables are created using keys. A key in one table will normally relate to a key in another table. Two tables in a database may also be unrelated.



# Database Relationships in MySQL

There are mainly 3 types of database relationships:

1. One-to-one (1:1) Relationship: If only one data in one table relates to the only one data in another table it is known as one-to-one (1:1) relationship.
2. One-to-many (1:M) Relationship: If only one data in one table relates to the multiple data in another table it is known as the one-to-many (1:M) relationship.
3. Many-to-many (M:M) Relationship: And if multiple data in one table relates to the multiple data in another table it is known as many-to-many (M:M) relationship.

# One-to-one (1:1) Relationship

In a One-to-One (1:1) Relationship only one data in one table relates to only one data in another table. Take a look at the example tables, we will use sakila database. We have two tables: customer and address. First table stores customer\_id, first\_name, last\_name, address\_id. The second table stores address\_id and address column values which are shown in the below image.

customer_id	store_id	first_name	last_name	email	address_id
1	1	MARY	SMITH	MARY.SMITH@sakilacustomer.org	2
2	1	PATRICIA	JOHNSON	PATRICIA.JOHNSON@sakilacustomer.org	3
3	1	LINDA	WILLIAMS	LINDA.WILLIAMS@sakilacustomer.org	1

address_id	address	address2	district	city_id	postal_code	phone
1	47 MySakila Drive	NULL	Alberta	300	NULL	14033335568
2	28 MySQL Boulevard	NULL	QLD	576	NULL	6172235589
3	23 Workhaven Lane	NULL	Alberta	300	NULL	14033335568

# One-to-one (1:1) Relationship

One-to-one relationships in databases are used for various reasons and serve specific purposes. Here are a few reasons why one might choose to use a one-to-one relationship in a database:

1. Security and access control: In some cases, you may want to restrict access to certain sensitive information. By placing sensitive data in a separate table with a one-to-one relationship, you can apply stricter access controls to that table, ensuring that only authorized users can view or modify the data.
2. Performance optimization: One-to-one relationships can help improve query performance. When you have large tables with a mix of frequently and infrequently accessed columns, splitting them into separate tables can reduce the overall data volume and improve query execution time.
3. Flexibility and scalability: One-to-one relationships provide flexibility and scalability in database design. If you anticipate that certain attributes may change or expand in the future, having separate tables allows you to easily add or modify columns without affecting the structure of the main table.




In this example, one customer can only have one address . and vice versa.

Look at the cardinality :

- || means one ( one customer only had one address )

# One-To-Many (1:M) Database Relationship

One-to-many relationship is a type of relationship between two entities in a database. It's called "one-to-many" because it describes a relationship where one entity (let's call it Entity A) can be associated with multiple instances of another entity (Entity B), but each instance of Entity B is associated with only one instance of Entity A.



In this example, one author can have many books , one book OR not have book at all. but one book only can have one author.

Look at the cardinality :

- || means one ( one book only had one author )
- o means zero ( one author can have 0 book )
- And <- means many ( one author can have many book )

# One-To-Many (1:M) Database Relationship

Take a look at the example between the "customers" and "rentals" tables. Each customer can have multiple rentals, but each rental belongs to only one customer. Here's an example SQL query that retrieves the rentals for a specific customer:




```
SELECT c.customer_id, r.rental_id, r.rental_date, r.return_date  
FROM rental r  
JOIN customer c ON r.customer_id = c.customer_id  
WHERE c.customer_id = 130;
```

customer_id	rental_id	rental_date	return_date
130	1	2005-05-24 22:53:30	2005-05-26 22:04:30
130	746	2005-05-29 09:25:10	2005-06-02 04:20:10
130	1630	2005-06-16 07:55:01	2005-06-19 06:38:01
130	1864	2005-06-17 01:39:47	2005-06-24 19:39:47

# Many-To-Many (M:M) Database Relationship

Many-to-many relationship is a type of relationship between two entities where each entity can be associated with multiple instances of the other entity. In this type of relationship, the cardinality is "many" on both sides.



Look at the relation, the cardinality is many in both side. In DBMS, we cannot implemented this relation of two tables. The solution is to made another table as composite table




Table film\_actors is the composite table, means is combination of multiple column from another table or combination of foreign key

# Many-To-Many (M:M) Database Relationship

Let's consider the Sakila database and demonstrate a many-to-many relationship using the "film" and "actor" tables. Each film can have multiple actors, and each actor can participate in multiple films. Here's an example query:

```
SELECT f.film_id, f.title, a.actor_id, a.first_name, a.last_name  
FROM film_actor fa  
JOIN film f ON fa.film_id = f.film_id  
JOIN actor a ON fa.actor_id = a.actor_id  
WHERE f.title = 'ACADEMY DINOSAUR';
```

# Many-To-Many (M:M) Database Relationship

In this example, we have the following tables involved:

- The "film" table contains film information, including the film\_id and title.
- The "actor" table contains actor information, including the actor\_id, first\_name, and last\_name.
- The "film\_actor" table serves as the junction table, connecting films and actors by storing the film\_id and actor\_id as foreign keys.
- The query retrieves the film\_id, title, actor\_id, first\_name, and last\_name for the film with the title 'ACADEMY DINOSAUR. By joining the "film\_actor" table with the "film" and "actor" tables, we establish the many-to-many relationship between films and actors.

```
SELECT f.film_id, f.title, a.actor_id, a.first_name, a.last_name
FROM film_actor fa
JOIN film f ON fa.film_id = f.film_id
JOIN actor a ON fa.actor_id = a.actor_id
WHERE f.title = 'ACADEMY DINOSAUR';
```

# Aggregate Function

An aggregate function in SQL performs a calculation on multiple values and returns a single value. SQL provides many aggregate functions that include avg, count, sum, min, max, etc. An aggregate function ignores NULL values when it performs the calculation, except for the count function.


An aggregate function in SQL returns one value after calculating multiple values of a column. We often use aggregate functions with the GROUP BY and HAVING clauses of the SELECT statement.

Various types of SQL aggregate functions are:

- Count()
- Sum()
- Avg()
- Min()
- Max()

# Aggregate Function - Count

The COUNT() function returns the number of rows in a database table



# Aggregate Function - Count

In this example, we have the following tables involved:

- The "category" table contains category information, including the category\_id and name.
- The "film\_category" table serves as the junction table, connecting categories and films by storing the category\_id and film\_id as foreign keys.
- The "film" table contains film information, including the film\_id.

The query uses the COUNT() function to count the number of films (film\_id) in each category. By joining the "category," "film\_category," and "film" tables, we establish the necessary relationships.



```
SELECT c.category_id, c.name AS category_name, COUNT(f.film_id) as film_count
FROM category c
JOIN film_category fc ON c.category_id = fc.category_id
JOIN film f ON fc.film_id = f.film_id
GROUP BY c.category_id, c.name;
```


# Aggregate Function - Count

The result set will include the category\_id, category\_name, and film\_count columns, where film\_count represents the count of films in each category. The GROUP BY clause is used to group the results by category\_id and category\_name.

category_id	category_name	film_count
1	Action	64
2	Animation	66
3	Children	60
4	Classics	57
5	Comedy	58
6	Documentary	68

# Aggregate Function - Sum

The SUM() function returns the total sum of a numeric column.



# Aggregate Function - Sum

Here's an example of using the SUM syntax in the Sakila database to calculate the total rental revenue for each customer:

- The "customer" table contains customer information, including the customer\_id, first\_name, and last\_name.
- The "payment" table contains payment information, including the customer\_id and amount.
- The query uses the SUM() function to calculate the total rental revenue (amount) for each customer. By joining the "customer" and "payment" tables based on the customer\_id, we establish the necessary relationship.

```
SELECT c.customer_id, c.first_name, c.last_name, SUM(p.amount) AS total_revenue
FROM customer c
JOIN payment p ON c.customer_id = p.customer_id
GROUP BY c.customer_id, c.first_name, c.last_name;
```

# Aggregate Function - Sum


The result set will include the `customer_id`, `first_name`, `last_name`, and `total_revenue` columns, where `total_revenue` represents the sum of payment amounts for each customer. The `GROUP BY` clause is used to group the results by `customer_id`, `first_name`, and `last_name`.

This query allows you to retrieve the total revenue generated from rentals for each customer in the Sakila database.

customer_id	first_name	last_name	total_revenue
1	MARY	SMITH	118.68
2	PATRICIA	JOHNSON	128.73
3	LINDA	WILLIAMS	135.74
4	BARBARA	JONES	81.78
5	ELIZABETH	BROWN	144.62
6	JENNIFER	DAVIS	93.72

# Aggregate Function - Avg

The AVG() function calculates the average of a set of values.



# Aggregate Function - Avg

In this example, we only need the "film" table, which contains information about films, including the length of each film.

The query uses the AVG() function to calculate the average length of films in the "film" table.

```
● ● ●  
SELECT AVG(length) as average_length  
FROM film;
```

# Aggregate Function - Avg

The result set will include a single column, average\_length, which represents the average length of all films in the database.

This query provides a simple and straightforward example of calculating the average length of films in the Sakila database.

average_length
115.2720

# Aggregate Function - Min

The MIN() aggregate function returns the lowest value (minimum) in a set of non-NULL values.



MIN( )

```
MIN( [ALL|DISTINCT] expression )
```

# Aggregate Function - Min

In this example, we are working with the "film" table, which contains information about films, including the length of each film.

The query uses the MIN() function to retrieve the minimum length among all films in the database.

```
SELECT MIN(length) as min_length  
FROM film;
```

# Aggregate Function - Min

The result set will include a single column, minimum\_length, which represents the minimum length among all films.

This query provides a simple and straightforward example of retrieving the minimum length in the Sakila database.

minimum_length
46

# Aggregate Function - Max

The MAX() aggregate function returns the highest value (maximum) in a set of non-NULL values.



MAX(  
     $\star$ )

MAX( [ALL|DISTINCT] expression )

# Aggregate Function - Max

In this example, we are working with the "film" table, which contains information about films, including the length of each film.

The query uses the MAX() function to retrieve the maximum length among all films in the database.

```
● ● ●  
SELECT MAX(length) as max_length  
FROM film;
```

# Aggregate Function - Max

The result set will include a single column, `minimum_length`, which represents the minimum length among all films.

This query provides a simple and straightforward example of retrieving the minimum length in the Sakila database.

maximum_length
185

# Aggregate Function with Group By and Having - Count

Still using sakila db as a sample data. In this case, we would like to create a feature about promotion based on table customer, and payment. For those customers who made a transaction more than 20 times, we would give some promotions to the next payment orders made. The question is, how many customers made more than 20?



# Aggregate Function with Group By and Having - Count

Based on that situation, we can combine several clauses in MySQL to provide that data. Please take a look at this query:

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Count

We can combine **COUNT**, **GROUP BY**, and **HAVING** in order to provide the data.

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Count

First, put **COUNT** at the beginning on the select statement to show the field as the expected output. Customer first\_name and last\_name would be shown as additional information.

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Count

Second, since we need to know detail information about customers, we join the table from payment to customer table using **INNER JOIN** and **ON** with customer\_id as a foreign key in

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Count

The **GROUP BY** clause in SQL is used to group rows with similar values together based on one or more columns. It is typically used in combination with aggregate functions to perform calculations on groups of rows rather than individual rows. In this case we combine with **COUNT** aggregate clause

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Count

At the end of the line, we uses **HAVING** clause in order to give a conditional result. Using comparison operator is part of **HAVING** clause in this case, we use more than 20 as a condition.

```
● ● ●  
SELECT COUNT(payment.customer_id), customer.first_name, customer.last_name  
FROM payment  
INNER JOIN customer  
ON customer.customer_id = payment.customer_id  
GROUP BY payment.customer_id  
HAVING COUNT(customer.customer_id) > 20;
```

# Aggregate Function with Group By and Having - Sum

Still using sakila db as a sample data. Last time, we just created a feature about discount promotion based on table customer, and payment based on 20 transaction minimum who would get the promotions. In this case is a little bit different, what if the one who got the discount promotions is based on



# Aggregate Function with Group By and Having - Sum

Please take a look at this query:

- The "customer" table contains customer information, including the customer\_id, first\_name, and last\_name.
- The "payment" table contains payment information, including the amount and customer\_id.

```
SELECT c.customer_id, c.first_name, c.last_name, SUM(p.amount) as total_payment
FROM customer c
JOIN payment p ON c.customer_id = p.customer_id
GROUP BY c.customer_id, c.first_name, c.last_name
HAVING SUM(p.amount) > 200;
```

# Aggregate Function with Group By and Having - Sum

The query uses the SUM() function to calculate the total payment made by each customer. By joining the "customer" and "payment" tables based on the customer\_id, we establish the necessary relationship.

```
SELECT c.customer_id, c.first_name, c.last_name, SUM(p.amount) as total_payment
FROM customer c
JOIN payment p ON c.customer_id = p.customer_id
GROUP BY c.customer_id, c.first_name, c.last_name
HAVING SUM(p.amount) > 200;
```

# Aggregate Function with Group By and Having - Sum

The result set will include the customer\_id, first\_name, last\_name, and total\_payment columns, where total\_payment represents the sum of payments made by each customer. The GROUP BY clause is used to group the results by customer\_id, first\_name, and last\_name.

```
SELECT c.customer_id, c.first_name, c.last_name, SUM(p.amount) as total_payment
FROM customer c
JOIN payment p ON c.customer_id = p.customer_id
GROUP BY c.customer_id, c.first_name, c.last_name
HAVING SUM(p.amount) > 200;
```

# Aggregate Function with Group By and Having - Sum

The HAVING clause is then used to filter the results and only include customers who have made a total payment greater than 200. The SUM() function is used to calculate the sum of payments for each customer and compare it to the threshold value of 200.

This query allows you to find customers in the Sakila database who have made a total payment greater than 200.

```
SELECT c.customer_id, c.first_name, c.last_name, SUM(p.amount) as total_payment
FROM customer c
JOIN payment p ON c.customer_id = p.customer_id
GROUP BY c.customer_id, c.first_name, c.last_name
HAVING SUM(p.amount) > 200;
```

# Subquery

- A Subquery or Inner query or a Nested query is a query within another SQL query and embedded within the WHERE clause.
- A subquery is used to return data that will be used in the main query as a condition to further restrict the data to be retrieved.
- Subqueries can be used with the SELECT, INSERT, UPDATE, and DELETE statements along with the operators like =, <, >, >=, <=, IN, BETWEEN, etc.

# Subquery Example Clause

```
SELECT film_id, title, rental_duration  
FROM film  
WHERE rental_duration > (  
    SELECT AVG(rental_duration) FROM film  
);|
```

# Subquery Example Clause

In this example, we have the "film" table involved.

The subquery is used to calculate the average rental duration of all films. It is nested within the main query.

The main query then retrieves the film\_id, title, and rental\_duration columns from the "film" table for films where the rental duration is greater than the average rental duration obtained from the subquery.

This query allows you to retrieve film information for films that have a rental duration greater than the average rental duration.

```
...  
SELECT film_id, title, rental_duration  
FROM film  
WHERE rental_duration > (  
    SELECT AVG(rental_duration) FROM film  
);|
```

# SQL Transaction

A **transaction** is a unit of work that is performed against a database. Transactions are units or sequences of work accomplished in a logical order, whether in a manual fashion by a user or automatically by some sort of a database program.

A transaction is the propagation of one or more changes to the database. For example, if you are creating a record or updating a record or deleting a record from the table, then you are performing a transaction on that table. It is important to control these transactions to ensure the data integrity and to handle database errors.

Practically, you will club many SQL queries into a group and you will execute all of them together as a part of a transaction.


# SQL Transaction - Control

COMMIT – to save the changes.

ROLLBACK – to roll back the changes.

SAVEPOINT – creates points within the groups of transactions in which to ROLLBACK.

SET TRANSACTION – Places a name on a transaction.




# PostgreSQL Transaction Statements

PostgreSQL provides us with the following important statements to control transactions:

- To start a transaction, you use the BEGIN statement. This begins a new transaction. BEGIN WORK is an alias for BEGIN.
- To commit the current transaction and make its changes permanent, you use the COMMIT statement.
- To roll back the current transaction and cancel its changes, you use the ROLLBACK statement.
- Unlike MySQL, PostgreSQL has **\*\*auto-commit\*\*** mode always disabled by default. Therefore, every operation is implicitly wrapped in a transaction unless explicitly started with BEGIN. No separate SET autocommit statement is needed.

# PostgreSQL Transaction Example

We will use the `orders` and `orderDetails` table for the demonstration



# PostgreSQL Transaction Commit Example

The following illustrates the step of creating a new sales order:

- First, start a transaction by using the **BEGIN** statement.
- Next, retrieve the next sales order number using the **nextval()** function from the sequence tied to the **orderNumber** column of the orders table.
- Then, insert a new sales order into the orders table using this sales order number.
- After that, insert sales order items into the orderdetails table using the same sales order number.
- Finally, commit the transaction using the **COMMIT** statement.
- Optionally, select data from both the **orders** and **orderdetails** tables to verify the new sales order.

# PostgreSQL Transaction Commit Example

```
● ● ●
-- 1. Start a new transaction
BEGIN;

-- 2. Get the latest order number
SELECT nextval('orders_orderNumber_seq') INTO orderNumber;

-- 3. Insert a new order fro customer 145
INSERT INTO orders(orderNumber, orderDate, requiredDate, shippedDate, status, customerNumber)
VALUES (
    orderNumber,
    '2005-05-31',
    '2005-06-10',
    '2005-06-11',
    'In Process',
    145
);

-- 4. Insert order line items
INSERT INTO orderdetails(orderNumber, productCode, quantityOrdered, priceEach, orderLineNumber)
VALUES
    (orderNumber, 'S18_1749', 30, 136, 1),
    (orderNumber, 'S18_2248', 50, 55, 2);

-- 5. Commit the transaction
COMMIT;
```

# PostgreSQL Transaction Rollback Example

In PostgreSQL, we can implement a rollback within a transaction to undo any changes made during the transaction before committing. A rollback cancels all operations done after the BEGIN statement and restores the database to its state before the transaction started.

- **Start a Transaction:** You begin by starting a transaction using the BEGIN statement.
- **Perform SQL Operations:** Execute the necessary SQL queries (inserts, updates, deletes, etc.).
- **Issue a Rollback:** If any errors occur or if you want to discard changes, you can use the ROLLBACK statement to cancel all the changes made within the transaction.
- **Check Results:** Optionally, verify that no changes were applied.

# PostgreSQL Transaction Rollback Example

```
● ● ●

-- 1. Start a new transaction
BEGIN;

-- 2. Perform some operations
INSERT INTO orders(orderNumber, orderDate, customerNumber)
VALUES (
    nextval('orders_orderNumber_seq'),
    '2025-05-31',
    145
);

-- 3. Another operation, potentially causing an issue
UPDATE customers SET creditLimit = creditLimit - 500 WHERE customerNumber = 145;

-- 4. Rollback the transaction if an error occurs or to discard changes
ROLLBACK;
```

# SQL Indexing

The **CREATE INDEX** statement is used to create indexes in tables.

Indexes are used to retrieve data from the database more quickly than otherwise. The users cannot see the indexes, they are just used to speed up searches/queries.

# SQL Indexing - Syntax

## Create Index Syntax,

Creates an index on a table. Duplicate values are allowed:

```
CREATE INDEX index_name  
ON table_name (column1, column2, ...);|
```

# SQL Indexing - Syntax

## Create Unique Index Syntax,

Creates an unique index on a table. Duplicate values are allowed:

```
CREATE UNIQUE INDEX index_name  
ON table_name (column1, column2, ...);
```

# SQL Indexing - Example

```
CREATE INDEX idx_lastname  
ON Persons (LastName);  
  
CREATE INDEX idx_pname  
ON Person (LastName, FirstName);|
```

# Thank you

